

```
In [179... #Load all packages and libraries
import pandas as pd
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn import linear_model
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [181... #Load dataset
diabetes = datasets.load_diabetes()
```

```
In [183... diabetes
```

```
Out[183... {'data': array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.0025922
6,
        0.01990749, -0.01764613],
        [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
        -0.06833155, -0.09220405],
        [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
        0.00286131, -0.02593034],
        ...,
        [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
        -0.04688253,  0.01549073],
        [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
        0.04452873, -0.02593034],
        [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
        -0.00422151,  0.00306441]]),
'target': array([151.,  75., 141., 206., 135.,  97., 138.,  63., 110.,
310., 101.,
        69., 179., 185., 118., 171., 166., 144.,  97., 168.,  68.,  4
9.,
        68., 245., 184., 202., 137.,  85., 131., 283., 129.,  59., 34
1.,
        87.,  65., 102., 265., 276., 252.,  90., 100.,  55.,  61.,  9
2.,
        259.,  53., 190., 142.,  75., 142., 155., 225.,  59., 104., 18
2.,
        128.,  52.,  37., 170., 170.,  61., 144.,  52., 128.,  71., 16
3.,
        150.,  97., 160., 178.,  48., 270., 202., 111.,  85.,  42., 17
0.,
        200., 252., 113., 143.,  51.,  52., 210.,  65., 141.,  55., 13
4.,
        42., 111.,  98., 164.,  48.,  96.,  90., 162., 150., 279.,  9
2.,
        83., 128., 102., 302., 198.,  95.,  53., 134., 144., 232.,  8
1.,
        104.,  59., 246., 297., 258., 229., 275., 281., 179., 200., 20
0.,
        173., 180.,  84., 121., 161.,  99., 109., 115., 268., 274., 15
8.,
        107.,  83., 103., 272.,  85., 280., 336., 281., 118., 317., 23
```

5.,  
1.,  
4.,  
7.,  
7.,  
9.,  
5.,  
5.,  
0.,  
7.,  
1.,  
6.,  
6.,  
3.,  
9.,  
2.,  
9.,  
5.,  
6.,  
1.,  
1.,  
8.,  
9.,  
8.,  
2.,  
5.,  
0.,  
2.,  
220., 57.] ),

60., 174., 259., 178., 128., 96., 126., 288., 88., 292., 7  
197., 186., 25., 84., 96., 195., 53., 217., 172., 131., 21  
59., 70., 220., 268., 152., 47., 74., 295., 101., 151., 12  
237., 225., 81., 151., 107., 64., 138., 185., 265., 101., 13  
143., 141., 79., 292., 178., 91., 116., 86., 122., 72., 12  
142., 90., 158., 39., 196., 222., 277., 99., 196., 202., 15  
77., 191., 70., 73., 49., 65., 263., 248., 296., 214., 18  
78., 93., 252., 150., 77., 208., 77., 108., 160., 53., 22  
154., 259., 90., 246., 124., 67., 72., 257., 262., 275., 17  
71., 47., 187., 125., 78., 51., 258., 215., 303., 243., 9  
150., 310., 153., 346., 63., 89., 50., 39., 103., 308., 11  
145., 74., 45., 115., 264., 87., 202., 127., 182., 241., 6  
94., 283., 64., 102., 200., 265., 94., 230., 181., 156., 23  
60., 219., 80., 68., 332., 248., 84., 200., 55., 85., 8  
31., 129., 83., 275., 65., 198., 236., 253., 124., 44., 17  
114., 142., 109., 180., 144., 163., 147., 97., 220., 190., 10  
191., 122., 230., 242., 248., 249., 192., 131., 237., 78., 13  
244., 199., 270., 164., 72., 96., 306., 91., 214., 95., 21  
263., 178., 113., 200., 139., 139., 88., 148., 88., 243., 7  
77., 109., 272., 60., 54., 221., 90., 311., 281., 182., 32  
58., 262., 206., 233., 242., 123., 167., 63., 197., 71., 16  
140., 217., 121., 235., 245., 40., 52., 104., 132., 88., 6  
219., 72., 201., 110., 51., 277., 63., 118., 69., 273., 25  
43., 198., 242., 232., 175., 93., 168., 275., 293., 281., 7  
140., 189., 181., 209., 136., 261., 113., 131., 174., 257., 5  
84., 42., 146., 212., 233., 91., 111., 152., 120., 67., 31  
94., 183., 66., 173., 72., 49., 64., 48., 178., 104., 13

```

'frame': None,
'DESCR': '.. _diabetes_dataset:\n\nDiabetes dataset\n-----\n\nTen baseline variables, age, sex, body mass index, average blood\npressure, and six blood serum measurements were obtained for each of n=\n442 diabetes patients, as well as the response of interest, a\nquantitative measure of disease progression one year after baseline.\n\n**Data Set Characteristics:**\n\nNumber of Instances: 442\n\nNumber of Attribute\ns: First 10 columns are numeric predictive values\n\nTarget: Column 11\nis a quantitative measure of disease progression one year after baseline\n\nAttribute Information:\n    - age      age in years\n    - sex\n    - bmi      body mass index\n    - bp      average blood pressure\n    - s1      tc, total serum cholesterol\n    - s2      ldl, low-density lipoproteins\n    - s3      hdl, high-density lipoproteins\n    - s4      tc h, total cholesterol / HDL\n    - s5      ltg, possibly log of serum tri glycerides level\n    - s6      glu, blood sugar level\n\nNote: Each of these 10 feature variables have been mean centered and scaled by the standard deviation times the square root of `n_samples` (i.e. the sum of squares of each column totals 1).\n\nSource URL:\nhttps://www4.stat.ncsu.edu/~boos/var.select/diabetes.html\n\nFor more information see:\nBradley Efron, Trevor Hastie, Iain Johnstone and Robert Tibshirani (2004) "Least Angle Regression," Annals of Statistics (with discussion), 407-499.\n(https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf)\n',
'feature_names': ['age',
'sex',
'bmi',
'bp',
's1',
's2',
's3',
's4',
's5',
's6'],
'data_filename': 'diabetes_data_raw.csv.gz',
'target_filename': 'diabetes_target.csv.gz',
'data_module': 'sklearn.datasets.data'}

```

```

In [185... # Show feature names
print(diabetes.feature_names)

```

```

['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']

```

```

In [187... # Create Feature and Target variables
X = diabetes.data
y = diabetes.target
print("dimension of X:", X.shape)
print("dimension of y:", y.shape)

```

```

dimension of X: (442, 10)
dimension of y: (442,)

```

```

In [189... # Split data into testing & training
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2

```

```

In [191... # Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)

```

```
X_test_scaled = scaler.transform(X_test)
```

```
In [193... # Dimensions of training data and testing data
print("feature training data:", X_train.shape, "/ target training data:",
print("feature testing data:", X_test.shape, "/ target testing data:" , y
print("scaled feature training data:", X_train_scaled.shape)
print("scaled feature test data:", X_test_scaled.shape)
```

```
feature training data: (353, 10) / target training data: (353,)
feature testing data: (89, 10) / target testing data: (89,)
scaled feature training data: (353, 10)
scaled feature test data: (89, 10)
```

```
In [195... #BUILDING LINEAR REGRESSION MODEL
# 1. defining the model
model = linear_model.LinearRegression()
```

```
In [197... # 2. build training model
model.fit(X_train_scaled, y_train)
```

```
Out[197... ▼ LinearRegression ⓘ ⓘ
LinearRegression()
```

```
In [199... # Applying training model to testing dataset
Y_pred = model.predict(X_test_scaled)
print(Y_pred)
```

```
[158.82689128 232.78373635 87.28268709 128.95236713 67.66313843
126.14110688 219.14585502 180.36516868 240.47185892 146.29544617
218.07251394 227.55217003 234.06388096 156.16293495 52.44746772
103.96782622 79.18036892 163.03148115 120.28946439 134.50558252
183.04263076 186.44995146 141.97441238 163.88064257 175.35236569
136.16146038 169.88077624 266.04171109 123.79984609 286.29077474
180.62640457 144.61102647 172.07594755 79.99445059 112.84725673
110.49595288 131.45259736 120.21140734 224.35302993 60.65423612
166.51559644 198.70405821 81.47531693 120.5324054 198.83101816
211.49369722 152.48611182 253.82678073 183.78230134 189.40414745
169.22566031 151.17807948 166.9624241 164.53463363 172.14749675
227.36590659 107.39736761 70.91704662 83.66134184 168.75349799
131.92010357 145.88116863 178.11036137 129.13969108 77.21187989
252.63863297 165.46166855 183.18564755 131.55842155 220.76088107
159.7281096 57.15148271 132.12656623 211.1858594 189.50235046
155.09555932 180.57209118 204.17449123 164.20392112 139.46522389
88.3872361 66.50617918 162.8997838 125.95341377 93.95661619
204.27347681 92.49964795 239.03740687 81.38544078]
```

```
In [201... # PREDICTION RESULTS
# PRINT MODEL PERFORMANCE
print("Parameters:", model.coef_) #coefficients of feature varibale
print("Intercept:", model.intercept_)
print("MSE: %.2f"
      % mean_squared_error(y_test, Y_pred))
print("Coefficient of determination: %.2f"
```

```
% r2_score(y_test, Y_pred))
```

```
Parameters: [ 2.68557884 -9.93762507 19.923154 15.19270396 -41.56270731
```

```
26.38957583 4.12720436 6.41289538 37.92535719 5.31541486]
```

```
Intercept: 153.44759206798867
```

```
MSE: 3410.87
```

```
Coefficient of determination:0.49
```

```
In [203... feature_names = diabetes.feature_names
print(feature_names[0])
```

age

```
In [97]: # Each model coefficient is multiplied by its corresponding feature varia
# The sum of all products of model coef and features should produce the v
#  $Y^{\wedge} = \text{model.coef}_{[0]} * \text{feature\_names}[0] + \text{model.coef}_{[1]} * \text{feature\_names}[1]$ 
#  $\text{model.coef}_{[10]} * \text{feature\_names}[10] = \text{model.intercept\_}$ 
```

```
In [205... y_test
```

```
Out[205... array([200., 321., 53., 89., 158., 71., 257., 262., 306., 116., 237.,
268., 246., 109., 48., 118., 59., 237., 64., 187., 144., 217.,
302., 154., 85., 42., 58., 220., 78., 270., 107., 88., 174.,
48., 71., 142., 83., 150., 248., 70., 131., 292., 37., 64.,
129., 202., 200., 308., 90., 142., 216., 129., 151., 168., 252.,
217., 68., 134., 201., 147., 71., 25., 70., 51., 77., 346.,
200., 167., 65., 180., 104., 39., 84., 249., 124., 81., 91.,
121., 259., 168., 55., 75., 155., 53., 64., 265., 115., 215.,
72.]])
```


```
In [207... Y_pred
```

```
Out[207... array([158.82689128, 232.78373635, 87.28268709, 128.95236713,
67.66313843, 126.14110688, 219.14585502, 180.36516868,
240.47185892, 146.29544617, 218.07251394, 227.55217003,
234.06388096, 156.16293495, 52.44746772, 103.96782622,
79.18036892, 163.03148115, 120.28946439, 134.50558252,
183.04263076, 186.44995146, 141.97441238, 163.88064257,
175.35236569, 136.16146038, 169.88077624, 266.04171109,
123.79984609, 286.29077474, 180.62640457, 144.61102647,
172.07594755, 79.99445059, 112.84725673, 110.49595288,
131.45259736, 120.21140734, 224.35302993, 60.65423612,
166.51559644, 198.70405821, 81.47531693, 120.5324054 ,
198.83101816, 211.49369722, 152.48611182, 253.82678073,
183.78230134, 189.40414745, 169.22566031, 151.17807948,
166.9624241 , 164.53463363, 172.14749675, 227.36590659,
107.39736761, 70.91704662, 83.66134184, 168.75349799,
131.92010357, 145.88116863, 178.11036137, 129.13969108,
77.21187989, 252.63863297, 165.46166855, 183.18564755,
131.55842155, 220.76088107, 159.7281096 , 57.15148271,
132.12656623, 211.1858594 , 189.50235046, 155.09555932,
180.57209118, 204.17449123, 164.20392112, 139.46522389,
88.3872361 , 66.50617918, 162.8997838 , 125.95341377,
93.95661619, 204.27347681, 92.49964795, 239.03740687,
81.38544078])
```


```
In [209... # VISUALIZATIONS - SCATTERPLOTS
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [211... # first look at data points on the plot
sns.scatterplot(x=y_test,y=Y_pred, )
```

Out[211... <Axes: >



```
In [213... # FINAL & MAIN SCATTER PLOT WITH REGRESSION/PREDICTION LINE
plt.figure(figsize = (12, 8))
plt.scatter(y_test, Y_pred, marker = '+', alpha =0.5, color='blue', label
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()],
         'r-',
         lw = 2,
         label = 'Prediction Line')
plt.title('Diabetes Prediction')
plt.xlabel('Actual Values (y_test)')
plt.ylabel('Prediction Values (Y_pred)')
plt.legend()
plt.grid(True)
plt.savefig("Diabetes Prediction(scaled)", dpi=300, bbox_inches = 'tight')
plt.show()
```



In [ ]: